

Chapter 8 Questions:

1. Run a few randomly-generated problems with just two jobs and two queues; compute the MLFQ execution trace for each. Make your life easier by limiting the length of each job and turning off I/Os.

Example1:

```
1e18@dhcp-172-17-32-41 cpu-sched-mlfq % ./mlfq.py -s 20 -n 2 -j 2 -i 0 -m 10 -M 0 -c
```

Here is the list of inputs:

OPTIONS jobs 2

OPTIONS queues 2

OPTIONS allotments for queue 1 is 1

OPTIONS quantum length for queue 1 is 10

OPTIONS allotments for queue 0 is 1

OPTIONS quantum length for queue 0 is 10

OPTIONS boost 0

OPTIONS ioTime 0

OPTIONS stayAfterIO False

OPTIONS iobump False

For each job, three defining characteristics are given:

startTime : at what time does the job enter the system

runTime : the total CPU time needed by the job to finish

ioFreq : every ioFreq time units, the job issues an I/O
(the I/O takes ioTime units to complete)

Job List:

Job 0: startTime 0 - runTime 9 - ioFreq 0

Job 1: startTime 0 - runTime 7 - ioFreq 0

Execution Trace:

[time 0] JOB BEGINS by JOB 0

[time 0] JOB BEGINS by JOB 1

[time 0] Run JOB 0 at PRIORITY 1 [TICKS 9 ALLOT 1 TIME 8 (of 9)]

[time 1] Run JOB 0 at PRIORITY 1 [TICKS 8 ALLOT 1 TIME 7 (of 9)]

[time 2] Run JOB 0 at PRIORITY 1 [TICKS 7 ALLOT 1 TIME 6 (of 9)]

[time 3] Run JOB 0 at PRIORITY 1 [TICKS 6 ALLOT 1 TIME 5 (of 9)]

[time 4] Run JOB 0 at PRIORITY 1 [TICKS 5 ALLOT 1 TIME 4 (of 9)]

[time 5] Run JOB 0 at PRIORITY 1 [TICKS 4 ALLOT 1 TIME 3 (of 9)]

[time 6] Run JOB 0 at PRIORITY 1 [TICKS 3 ALLOT 1 TIME 2 (of 9)]

[time 7] Run JOB 0 at PRIORITY 1 [TICKS 2 ALLOT 1 TIME 1 (of 9)]

[time 8] Run JOB 0 at PRIORITY 1 [TICKS 1 ALLOT 1 TIME 0 (of 9)]

```
[ time 9 ] FINISHED JOB 0
[ time 9 ] Run JOB 1 at PRIORITY 1 [ TICKS 9 ALLOT 1 TIME 6 (of 7) ]
[ time 10 ] Run JOB 1 at PRIORITY 1 [ TICKS 8 ALLOT 1 TIME 5 (of 7) ]
[ time 11 ] Run JOB 1 at PRIORITY 1 [ TICKS 7 ALLOT 1 TIME 4 (of 7) ]
[ time 12 ] Run JOB 1 at PRIORITY 1 [ TICKS 6 ALLOT 1 TIME 3 (of 7) ]
[ time 13 ] Run JOB 1 at PRIORITY 1 [ TICKS 5 ALLOT 1 TIME 2 (of 7) ]
[ time 14 ] Run JOB 1 at PRIORITY 1 [ TICKS 4 ALLOT 1 TIME 1 (of 7) ]
[ time 15 ] Run JOB 1 at PRIORITY 1 [ TICKS 3 ALLOT 1 TIME 0 (of 7) ]
[ time 16 ] FINISHED JOB 1
```

Final statistics:

Job 0: startTime 0 - response 0 - turnaround 9

Job 1: startTime 0 - response 9 - turnaround 16

Avg 1: startTime n/a - response 4.50 - turnaround 12.50

Example2:

```
1e18@dhcp-172-17-32-41 cpu-sched-mlfq % ./mlfq.py -s 20 -n 2 -j 2 -i 0 -m 20 -M 0 -c
```

Here is the list of inputs:

OPTIONS jobs 2

OPTIONS queues 2

OPTIONS allotments for queue 1 is 1

OPTIONS quantum length for queue 1 is 10

OPTIONS allotments for queue 0 is 1

OPTIONS quantum length for queue 0 is 10

OPTIONS boost 0

OPTIONS ioTime 0

OPTIONS stayAfterIO False

OPTIONS iobump False

For each job, three defining characteristics are given:

startTime : at what time does the job enter the system

runTime : the total CPU time needed by the job to finish

ioFreq : every ioFreq time units, the job issues an I/O
(the I/O takes ioTime units to complete)

Job List:

Job 0: startTime 0 - runTime 18 - ioFreq 0

Job 1: startTime 0 - runTime 15 - ioFreq 0

Execution Trace:

[time 0] JOB BEGINS by JOB 0
 [time 0] JOB BEGINS by JOB 1
 [time 0] Run JOB 0 at PRIORITY 1 [TICKS 9 ALLOT 1 TIME 17 (of 18)]
 [time 1] Run JOB 0 at PRIORITY 1 [TICKS 8 ALLOT 1 TIME 16 (of 18)]
 [time 2] Run JOB 0 at PRIORITY 1 [TICKS 7 ALLOT 1 TIME 15 (of 18)]
 [time 3] Run JOB 0 at PRIORITY 1 [TICKS 6 ALLOT 1 TIME 14 (of 18)]
 [time 4] Run JOB 0 at PRIORITY 1 [TICKS 5 ALLOT 1 TIME 13 (of 18)]
 [time 5] Run JOB 0 at PRIORITY 1 [TICKS 4 ALLOT 1 TIME 12 (of 18)]
 [time 6] Run JOB 0 at PRIORITY 1 [TICKS 3 ALLOT 1 TIME 11 (of 18)]
 [time 7] Run JOB 0 at PRIORITY 1 [TICKS 2 ALLOT 1 TIME 10 (of 18)]
 [time 8] Run JOB 0 at PRIORITY 1 [TICKS 1 ALLOT 1 TIME 9 (of 18)]
 [time 9] Run JOB 0 at PRIORITY 1 [TICKS 0 ALLOT 1 TIME 8 (of 18)]
 [time 10] Run JOB 1 at PRIORITY 1 [TICKS 9 ALLOT 1 TIME 14 (of 15)]
 [time 11] Run JOB 1 at PRIORITY 1 [TICKS 8 ALLOT 1 TIME 13 (of 15)]
 [time 12] Run JOB 1 at PRIORITY 1 [TICKS 7 ALLOT 1 TIME 12 (of 15)]
 [time 13] Run JOB 1 at PRIORITY 1 [TICKS 6 ALLOT 1 TIME 11 (of 15)]
 [time 14] Run JOB 1 at PRIORITY 1 [TICKS 5 ALLOT 1 TIME 10 (of 15)]
 [time 15] Run JOB 1 at PRIORITY 1 [TICKS 4 ALLOT 1 TIME 9 (of 15)]
 [time 16] Run JOB 1 at PRIORITY 1 [TICKS 3 ALLOT 1 TIME 8 (of 15)]
 [time 17] Run JOB 1 at PRIORITY 1 [TICKS 2 ALLOT 1 TIME 7 (of 15)]
 [time 18] Run JOB 1 at PRIORITY 1 [TICKS 1 ALLOT 1 TIME 6 (of 15)]
 [time 19] Run JOB 1 at PRIORITY 1 [TICKS 0 ALLOT 1 TIME 5 (of 15)]
 [time 20] Run JOB 0 at PRIORITY 0 [TICKS 9 ALLOT 1 TIME 7 (of 18)]
 [time 21] Run JOB 0 at PRIORITY 0 [TICKS 8 ALLOT 1 TIME 6 (of 18)]
 [time 22] Run JOB 0 at PRIORITY 0 [TICKS 7 ALLOT 1 TIME 5 (of 18)]
 [time 23] Run JOB 0 at PRIORITY 0 [TICKS 6 ALLOT 1 TIME 4 (of 18)]
 [time 24] Run JOB 0 at PRIORITY 0 [TICKS 5 ALLOT 1 TIME 3 (of 18)]
 [time 25] Run JOB 0 at PRIORITY 0 [TICKS 4 ALLOT 1 TIME 2 (of 18)]
 [time 26] Run JOB 0 at PRIORITY 0 [TICKS 3 ALLOT 1 TIME 1 (of 18)]
 [time 27] Run JOB 0 at PRIORITY 0 [TICKS 2 ALLOT 1 TIME 0 (of 18)]
 [time 28] FINISHED JOB 0
 [time 28] Run JOB 1 at PRIORITY 0 [TICKS 9 ALLOT 1 TIME 4 (of 15)]
 [time 29] Run JOB 1 at PRIORITY 0 [TICKS 8 ALLOT 1 TIME 3 (of 15)]
 [time 30] Run JOB 1 at PRIORITY 0 [TICKS 7 ALLOT 1 TIME 2 (of 15)]
 [time 31] Run JOB 1 at PRIORITY 0 [TICKS 6 ALLOT 1 TIME 1 (of 15)]
 [time 32] Run JOB 1 at PRIORITY 0 [TICKS 5 ALLOT 1 TIME 0 (of 15)]
 [time 33] FINISHED JOB 1

Final statistics:

Job 0: startTime 0 - response 0 - turnaround 28

Job 1: startTime 0 - response 10 - turnaround 33

Avg 1: startTime n/a - response 5.00 - turnaround 30.50

Example 3:

```
1e18@dhcp-172-17-32-41 cpu-sched-mlfq % ./mlfq.py -s 20 -n 2 -j 2 -i 0 -m 5 -M 0 -c
```

Here is the list of inputs:

OPTIONS jobs 2

OPTIONS queues 2

OPTIONS allotments for queue 1 is 1

OPTIONS quantum length for queue 1 is 10

OPTIONS allotments for queue 0 is 1

OPTIONS quantum length for queue 0 is 10

OPTIONS boost 0

OPTIONS ioTime 0

OPTIONS stayAfterIO False

OPTIONS iobump False

For each job, three defining characteristics are given:

startTime : at what time does the job enter the system

runTime : the total CPU time needed by the job to finish

ioFreq : every ioFreq time units, the job issues an I/O
(the I/O takes ioTime units to complete)

Job List:

Job 0: startTime 0 - runTime 4 - ioFreq 0

Job 1: startTime 0 - runTime 4 - ioFreq 0

Execution Trace:

[time 0] JOB BEGINS by JOB 0

[time 0] JOB BEGINS by JOB 1

[time 0] Run JOB 0 at PRIORITY 1 [TICKS 9 ALLOT 1 TIME 3 (of 4)]

[time 1] Run JOB 0 at PRIORITY 1 [TICKS 8 ALLOT 1 TIME 2 (of 4)]

[time 2] Run JOB 0 at PRIORITY 1 [TICKS 7 ALLOT 1 TIME 1 (of 4)]

[time 3] Run JOB 0 at PRIORITY 1 [TICKS 6 ALLOT 1 TIME 0 (of 4)]

[time 4] FINISHED JOB 0

[time 4] Run JOB 1 at PRIORITY 1 [TICKS 9 ALLOT 1 TIME 3 (of 4)]

[time 5] Run JOB 1 at PRIORITY 1 [TICKS 8 ALLOT 1 TIME 2 (of 4)]

[time 6] Run JOB 1 at PRIORITY 1 [TICKS 7 ALLOT 1 TIME 1 (of 4)]

[time 7] Run JOB 1 at PRIORITY 1 [TICKS 6 ALLOT 1 TIME 0 (of 4)]

[time 8] FINISHED JOB 1

Final statistics:

Job 0: startTime 0 - response 0 - turnaround 4

Job 1: startTime 0 - response 4 - turnaround 8

Avg 1: startTime n/a - response 2.00 - turnaround 6.00

2. How would you run the scheduler to reproduce each of the examples in the chapter?

Example1: A Single Long-Running job

`./mlfq.py -n 3 -q 10 -l 0,100,0 -c`

Example2: Along came a interactive job

`./mlfq.py -n 3 -q 10 -l 0,100,0:50,10,0 -c`

Example3: A Mixed I/O-intensive and CPU-intensive Workload

`./mlfq.py -n 3 -q 10 -l 0,50,0:10,15,1 -i 5 -S -c`

Example4: Without Priority Boost

`./mlfq.py -n 3 -q 10 -l 0,120,0:100,50,1:100,50,1 -i 1 -S -c`

Example5: With Priority Boost

`./mlfq.py -n 3 -q 10 -l 0,150,0:100,50,1:100,50,1 -i 2 -S -B 40 -c`

Example6: Without Gaming Tolerance

`./mlfq.py -n 3 -l 0,100,0:30,90,9 -i 1 -S -c`

Example7: With Game Tolerance

`./mlfq.py -n 3 -q 10 -l 0,100,0:30,90,9 -i 1 -c`

Example8: Lower Priority, Longer Quanta

`./mlfq.py -n 3 -Q 10,20,30 -l 0,200,0:0,200,0 -c`

3. How would you configure the scheduler parameters to behave just like a round-robin scheduler?

When there is only one queue, the jobs on one queue would be executed in time slices.

4. Craft a workload with two jobs and scheduler parameters so that one job takes advantage of the older Rules 4a and 4b (turned on with the -S flag) to game the scheduler and obtain 99% of the CPU over a particular time interval.

`./mlfq.py -n 3 -l 0,100,0:30,90,9 -i 1 -S -c`

5. Given a system with a quantum length of 10ms in its highest queue, how often would you have to boost jobs back to the highest priority level (with the -B flag) to guarantee that a single longrunning (and potentially starving) job gets at least 5% of the CPU?

$T(\text{boost}) = 10 / 5\% = 200.$

For each turn of the long job, there are 10ms to run. It's at least 5% CPU. Thus $10/5\% = 200$ which means every 200 ms the long job needs to get turn to run.

6. One question that arises in scheduling is which end of a queue to add a job that just finished I/O; the -l flag changes this behavior for this scheduling simulator. Play around with some workloads and see if you can see the effect of this flag.

-l means that if specified, jobs that finished I/O move immediately to the front of the current queue.

```
./mlfq.py -n 2 -q 10 -l 0,30,0:0,30,11 -i 1 -S -c
```

```
./mlfq.py -n 2 -q 10 -l 0,30,0:0,30,11 -i 1 -S -l -c
```

Chapter 9 Questions:

1. Compute the solutions for simulations with 3 jobs and random seeds of 1, 2, and 3.

```
./lottery.py -j 3 -s 1 -c
```

```
./lottery.py -j 3 -s 2 -c
```

```
./lottery.py -j 3 -s 3 -c
```

2. Now run with two specific jobs: each of length 10, but one (job 0) with just 1 ticket and the other (job 1) with 100 (e.g., -l 10:1,10:100). What happens when the number of tickets is so imbalanced? Will job 0 ever run before job 1 completes? How often? In general, what does such a ticket imbalance do to the behavior of lottery scheduling?

```
./lottery.py -l 10:1,10:100 -c
```

Job0 would rarely have chance to run before job1 completes. The probability is 1/100.

A ticket imbalance would make lottery scheduling meaningless.

3. When running with two jobs of length 100 and equal ticket allocations of 100 (-l 100:100,100:100), how unfair is the scheduler? Run with some different random seeds to determine the (probabilistic) answer; let unfairness be determined by how much earlier one job finishes than the other.

```
./lottery.py -s 0 -l 100:100,100:100 -c
```

$U = 192/200 = 0.96$

```
./lottery.py -s 1 -l 100:100,100:100 -c
```

$U = 196/200 = 0.98$

```
./lottery.py -s 2 -l 100:100,100:100 -c
```

$U = 190/200 = 0.95$

```
./lottery.py -s 3 -l 100:100,100:100 -c
```

$U = 196/200 = 0.98$

4. How does your answer to the previous question change as the quantum size (-q) gets larger?

```
./lottery.py -q 10 -s 0 -l 100:100,100:100 -c
```

$U = 150/200 = 0.75$

```
./lottery.py -q 10 -s 1 -l 100:100,100:100 -c
```

$U = 160/200 = 0.8$

```
./lottery.py -q 10 -s 2 -l 100:100,100:100 -c
```

$U = 190/200 = 0.95$

```
./lottery.py -q 10 -s 3 -l 100:100,100:100 -c
```

$U = 190/200 = 0.95$

```
./lottery.py -q 50 -s 0 -l 100:100,100:100 -c
```

$U = 150/200 = 0.75$

```
./lottery.py -q 50 -s 1 -l 100:100,100:100 -c
```

$U = 150/200 = 0.75$

```
./lottery.py -q 50 -s 2 -l 100:100,100:100 -c
```

$U = 100/200 = 0.5$

```
./lottery.py -q 50 -s 3 -l 100:100,100:100 -c
```

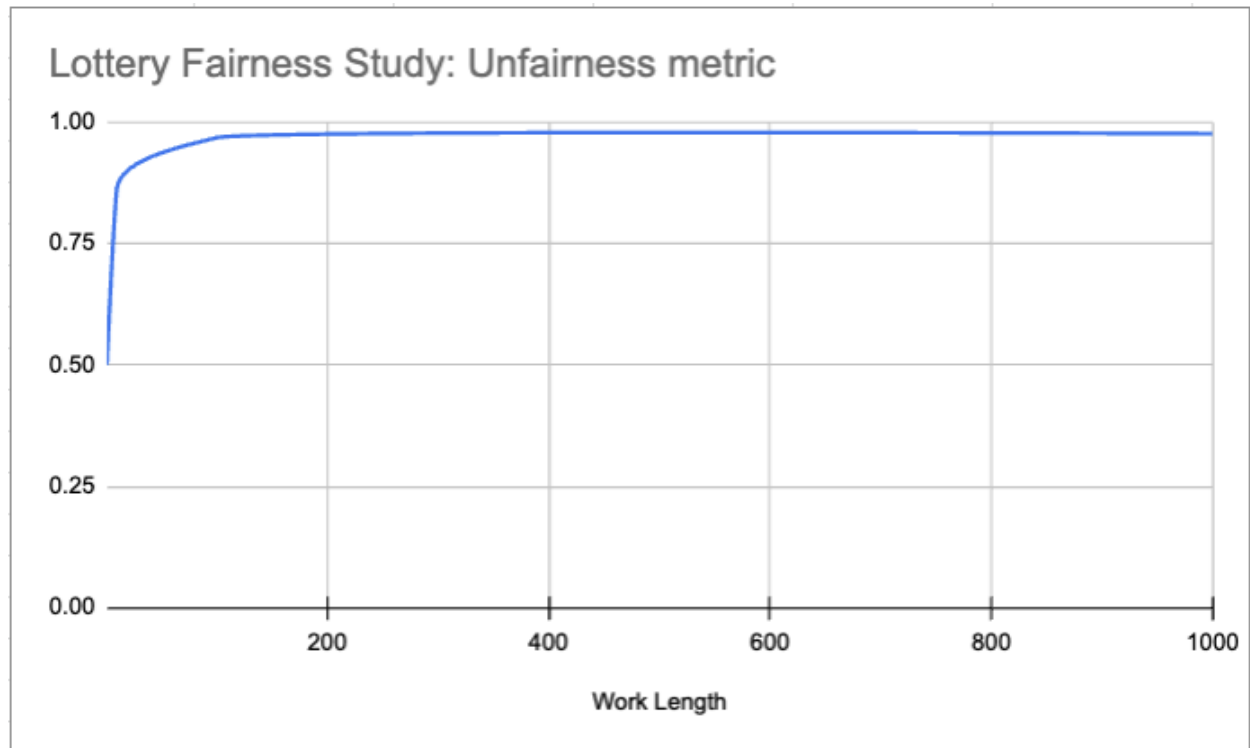
$U = 150/200 = 0.75$

When the quantum size gets larger, unfairness metric turns to be smaller which means it's more unfair than before.

5. Can you make a version of the graph that is found in the chapter? What else would be worth exploring? How would the graph look with a stride scheduler?

Work length:

Work Length	1	10	100	1000
U(seed = 0)	0.5	0.75	0.96	0.971
U(seed = 1)	0.5	0.8	0.98	0.9515
U(seed = 2)	0.5	0.95	0.95	0.9835
U(seed = 3)	0.5	0.95	0.98	0.997
U(average)	0.5	0.8625	0.9675	0.97575



When we are using stride scheduler, the graph should be a horizontal line because the stride scheduler has the precision to get each job exactly right at the end of each scheduling cycle.